Manisha Sankalamaddi

United States • manishareddy3217@gmail.com • 816-263-5480 • LinkedIn • manishareddy.com

OBJECTIVE

Snowflake SnowPro Core Certified data professional with 3+ years of experience designing data pipelines, building cloud-based data warehouse solutions, and developing advanced SQL transformations for scalable analytics environments. Strong background in Snowflake, advanced SQL, stored procedures, ETL/ELT workflows, data migration, and cloud data engineering across Snowflake, AWS, and big data platforms. Experienced in building incremental data pipelines, curated reporting layers, and performance-optimized datasets that support analytics, reporting, and business decision-making. Seeking a Snowflake Developer role where I can contribute to data pipeline development, query optimization, data warehousing, and cloud-based data transformation initiatives.

TECHNICAL SKILLS

Snowflake: Snowflake Cloud Data Platform, Snowflake SQL, Stored Procedures, Streams, Tasks, External Stages, Zero-Copy Cloning, Incremental Loads, Query Optimization, Performance Tuning
Data Engineering / Warehousing: ETL/ELT, Data Migration, Data Transformation, Data Modeling, Dimensional Modeling, Star Schema, Data Warehousing, Data Validation, Data Reconciliation, Data Quality
Cloud Platforms: AWS S3, AWS Glue, Amazon Redshift, Amazon Athena, GCP BigQuery
Big Data: Hadoop, HDFS, Apache Spark, PySpark, Spark SQL
Programming: SQL, Python, Pandas, NumPy
BI / Reporting: Power BI, Tableau, KPI Reporting, Dashboard Development
DevOps / Tools: Azure DevOps, Git, CI/CD, Docker, Kubernetes

EXPERIENCE

Data Analyst – Snowflake Developer | Integrated Proteins *Oct 2025 – Present*

- Designed and developed Snowflake-based ELT pipelines to migrate operational data from source systems such as NewBoss and EveryWeigh into external stages, staging tables, and curated warehouse tables.
- Built and maintained advanced Snowflake SQL stored procedures for data ingestion, transformation, business-rule processing, and recurring load execution across reporting datasets.
- Implemented incremental load workflows using Snowflake Streams and Tasks, enabling efficient processing of newly added and changed records while reducing manual refresh effort.
- Automated recurring data workflows through Snowflake Tasks with CRON scheduling, ensuring timely delivery of refreshed datasets for analytics and reporting use cases.
- Created reporting-ready Snowflake views and transformed data layers used as the backend for Power BI dashboards, supporting sales, operations, and business performance reporting.
- Used Zero-Copy Cloning for backup, testing, and environment replication, helping preserve production data integrity while supporting development and validation activities.
- Performed data validation, reconciliation, deduplication, and query optimization in Snowflake to improve data quality, warehouse performance, and reliability of downstream business reporting.

Data Analyst | Sodexo *Jul 2024 – Oct 2025*

- Built and maintained AWS-based data migration and transformation workflows using Amazon S3, AWS Glue, Amazon Redshift, and Athena to support analytics and reporting requirements.
- Developed ETL pipelines to ingest raw data, apply transformations, and load structured datasets into Amazon Redshift for downstream analysis and BI consumption.
- Wrote advanced SQL queries in Redshift and Athena for transformation, reconciliation, validation, and trend analysis across operational and business datasets.
- Used Python (Pandas, NumPy) to clean, standardize, preprocess, and automate recurring data preparation tasks for reporting workflows.
- Created and maintained curated reporting tables and analytical datasets that improved consistency of KPIs and usability of downstream reporting assets.
- Supported incremental and scheduled data refresh workflows, ensuring timely movement of updated data from source layers into analytics-ready structures.
- Applied strong data quality controls, including source-to-target validation, reconciliation, deduplication, and issue resolution, while also supporting Power BI reporting built on curated AWS-backed datasets.

Data Engineer | RSoft Systems & Services Pvt. Ltd. *Jun 2021 – Jun 2023*

- Ingested data from external sources such as databases, flat files, APIs, and system-generated files into Hadoop/HDFS for centralized storage and large-scale distributed processing.
- Built ingestion workflows using ETL jobs and Python scripts to land source data into raw HDFS layers and prepare it for downstream transformation.
- Used PySpark and Spark SQL to read data from HDFS, clean and standardize fields, apply business rules, and transform raw data into structured analytics-ready datasets.
- Processed both historical and incremental datasets in Hadoop and Spark, converting high-volume raw data into curated outputs for reporting and downstream consumption.
- Wrote Spark SQL and SQL-based logic for joins, aggregations, validation, reconciliation, and deduplication across large-scale datasets.
- Moved transformed datasets from the Hadoop/Spark layer into AWS, including Amazon S3, to support cloud-based storage and analytics access.
- Prepared curated datasets for downstream Power BI and Tableau reporting, while maintaining strong data quality, consistency, and transformation accuracy across the pipeline.

EDUCATION

Master of Science, Computer Science *Aug 2023 – May 2025*
University of Missouri-Kansas City • GPA: 3.74 / 4.0

CERTIFICATIONS

Snowflake SnowPro Core Certified (Credential ID: 179071720)

Manisha Sankalamaddi - Projects

github.com/smanisha-reddy

PROJECTS

smanisha-reddy/Taxi_HDFS_To_Snowflake

Built an end-to-end NYC Taxi data pipeline using HDFS, PySpark, Snowflake, S3, and Streamlit. Ingested raw trip data, transformed it into analytics-ready datasets, loaded reference data from S3, created reporting views, and used Snowflake Cortex Code to quickly debug and resolve a Spark-to-Snowflake authentication issue.

smanisha-reddy/aws-s3-snowflake-products-pipeline

AWS S3 to Snowflake data pipeline for product data ingestion, transformation, and task-based automation using Snowflake storage integrations, external stages, raw staging tables, curated analytics tables, and scheduled refresh tasks.

smanisha-reddy/github-events-snowflake-pipeline

A free-access GitHub Events to Snowflake pipeline built with Python and Snowflake. Fetches GitHub event data through authenticated REST API calls, lands the data as NDJSON, loads it into Snowflake RAW tables, and transforms it into CURATED tables for analysis.

smanisha-reddy/streamlit_cortex_dashboard

Interactive Streamlit sales analytics dashboard built with Python and Snowflake, enhanced using Cortex Code CLI for prompt-driven code updates, UI improvements, and development acceleration.